

Database Systems




Department of Computer Science
Chungbuk National University

Prof. Jong Yun LEE

CHAPTER 7:

More SQL: Complex Queries, Triggers, Views, and Schema Modification



Chapter 7 Outline

- More Complex SQL Retrieval Queries
- Specifying Semantic Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Modification in SQL

More Complex SQL Retrieval Queries

More Complex SQL Retrieval Queries

- Additional features allow users to specify more complex retrievals from database:
 - Nested queries, joined tables, and outer joins (in the FROM clause), aggregate functions, and grouping

Comparisons Involving NULL and Three-Valued Logic

- Meanings of **NULL**
 - **Unknown value**
 - **Unavailable or withheld value**
 - **Not applicable attribute**
- Each individual NULL value considered to be different from every other NULL value
- **SQL uses a three-valued logic:**
 - TRUE, FALSE, and UNKNOWN (like Maybe)
- **NULL = NULL comparison is avoided**

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

Table 7.1 Logical Connectives in Three-Valued Logic

(a)	AND	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	OR	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	NOT			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

- **SQL allows queries that check whether an attribute value is NULL**
 - **IS** or **IS NOT NULL**

Query 18. Retrieve the names of all employees who do not have supervisors.

```
Q18:  SELECT  Fname, Lname
      FROM    EMPLOYEE
      WHERE   Super_ssn IS NULL;
```


Comparisons Involving NULL and Three-Valued Logic (cont'd.)

```
mysql> desc current_dept_emp;
```

Field	Type	Null	Key	Default	Extra
emp_no	int(11)	NO		NULL	
dept_no	char(4)	NO		NULL	
from_date	date	YES		NULL	
to_date	date	YES		NULL	

```
4 rows in set (0.00 sec)
```

```
mysql> SELECT emp_no from current_dept_emp WHERE from_date  
IS NULL;
```

```
Empty set (0.33 sec)
```

```
mysql>
```

Nested Queries, Tuples, and Set/Multiset Comparisons

- **Nested queries**
 - Complete [select-from-where blocks within WHERE](#) clause of another query
 - **Outer query and nested subqueries**
- **Comparison operator IN**
 - Compares value v with a set (or multiset) of values V
 - Evaluates to TRUE if v is one of the elements in V

Nested Queries (cont'd.)

```
Q4A:  SELECT  DISTINCT Pnumber
      FROM    PROJECT
      WHERE   Pnumber IN
            ( SELECT  Pnumber
              FROM    PROJECT, DEPARTMENT, EMPLOYEE
              WHERE   Dnum = Dnumber AND
                    Mgr_ssn = Ssn AND Lname = 'Smith' )

      OR

      Pnumber IN
      ( SELECT  Pno
        FROM    WORKS_ON, EMPLOYEE
        WHERE   Essn = Ssn AND Lname = 'Smith' );
```

Nested Queries (cont'd.)

```
mysql> SELECT emp_no FROM dept_emp WHERE emp_no IN  
( SELECT emp_no FROM current_dept_emp WHERE dept_no  
= 'd001') limit 5;
```

```
+-----+  
| emp_no |  
+-----+  
|  10017 |  
|  10055 |  
|  10058 |  
|  10108 |  
|  10108 |  
+-----+
```

```
5 rows in set (0.29 sec)
```

```
mysql>
```

Nested Queries (cont'd.)

- **Use tuples** of values in comparisons
 - Place them within parentheses

```
SELECT    DISTINCT Essn
FROM      WORKS_ON
WHERE     (Pno, Hours) IN ( SELECT    Pno, Hours
                           FROM      WORKS_ON
                           WHERE     Essn = '123456789' );
```

```
mysql> SELECT emp_no FROM dept_emp WHERE (emp_no, dept_no) IN
(SELECT emp_no, dept_no FROM dept_emp WHERE dept_no = 'd001')
limit 3;
```

```
+-----+
| emp_no |
+-----+
|  10017 |
|  10055 |
|  10058 |
+-----+
```

```
3 rows in set (0.00 sec)
```

```
mysql>
```

Nested Queries (cont'd.)

- **Use other comparison operators to compare a single value v**
 - **= ANY (or = SOME) operator**
 - Returns TRUE if the value v is equal to some value in the set V and is hence equivalent to IN
 - Other operators that can be combined with ANY (or SOME): $>$, $>=$, $<$, $<=$, and $<>$
 - ALL: value must exceed all values from nested query

```
SELECT      Lname, Fname
FROM        EMPLOYEE
WHERE       Salary > ALL      ( SELECT      Salary
                                FROM        EMPLOYEE
                                WHERE       Dno = 5 );
```

Nested Queries (cont'd.)

```
mysql> desc current_dept_emp;
```

Field	Type	Null	Key	Default	Extra
emp_no	int(11)	NO		NULL	
dept_no	char(4)	NO		NULL	
from_date	date	YES		NULL	
to_date	date	YES		NULL	

```
4 rows in set (0.00 sec)
```

```
mysql> desc salaries;
```

Field	Type	Null	Key	Default	Extra
emp_no	int(11)	NO	PRI	NULL	
salary	int(11)	NO		NULL	
from_date	date	NO	PRI	NULL	
to_date	date	NO		NULL	

```
4 rows in set (0.00 sec)
```

```
mysql> SELECT salary FROM salaries WHERE salary > ALL (SELECT S.salary from salaries S,  
current_dept_emp C WHERE S.emp_no = C.emp_no AND C.dept_no = 'd001') limit 3;
```

salary
145732
145215
148820

Nested Queries (cont'd.)

- **Avoid potential errors and ambiguities**
 - Create tuple variables ([aliases](#)) for all tables referenced in SQL query

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT      E.Fname, E.Lname
      FROM        EMPLOYEE AS E
      WHERE       E.Ssn IN  ( SELECT      D.Essn
                              FROM        DEPENDENT AS D
                              WHERE       E.Fname = D.Dependent_name
                              AND E.Sex = D.Sex );
```


Nested Queries (cont'd.)

```
mysql> SELECT salary FROM salaries WHERE salary > ALL (SELECT S.salary
from salaries AS S, current_dept_emp AS C WHERE S.emp_no = C.emp_no
AND C.dept_no = 'd001') limit 5;
```

```
+-----+
| salary |
+-----+
| 145732  |
| 145215  |
| 148820  |
| 145300  |
| 149440  |
+-----+
```

```
5 rows in set (0.84 sec)
```

```
mysql>
```

Correlated Nested Queries

- Queries that are nested using the = or IN comparison operator can be collapsed into one single block: E.g., Q16 can be written as:

```
Q16:  SELECT  E.Fname, E.Lname
      FROM    EMPLOYEE AS E
      WHERE   E.Ssn IN ( SELECT  D.Essn
                        FROM    DEPENDENT AS D
                        WHERE   E.Fname = D.Dependent_name
                        AND E.Sex = D.Sex );
```

```
Q16A:  SELECT  E.Fname, E.Lname
      FROM    EMPLOYEE AS E, DEPENDENT AS D
      WHERE   E.Ssn=D.Essn AND E.Sex=D.Sex
      AND E.Fname=D.Dependent_name;
```

- **Correlated nested query**
 - Evaluated once for each tuple in the outer query

The EXISTS and UNIQUE Functions in SQL for correlating queries

- **EXISTS** function
 - Check whether the result of a correlated nested query is empty or not. They are Boolean functions that return a TRUE or FALSE result.
- **EXISTS and NOT EXISTS**
 - Typically used in conjunction with a correlated nested query
- **SQL function UNIQUE(Q)**
 - Returns TRUE if there are no duplicate tuples in the result of query Q

USE of EXISTS

Q7: List the names of managers who have at least one dependent.

```
SELECT Fname, Lname
FROM Employee
WHERE EXISTS (SELECT *
               FROM DEPENDENT
               WHERE Ssn= Essn) AND
EXISTS (SELECT *
        FROM Department
        WHERE Ssn= Mgr_Ssn)
```

USE of EXISTS

```
mysql> SELECT salary FROM salaries WHERE EXISTS (SELECT S.salary from
salaries AS S, current_dept_emp AS C WHERE S.emp_no = C.emp_no AND
C.dept_no = 'd001') limit 5;
```

```
+-----+
| salary |
+-----+
|  60117 |
|  62102 |
|  66074 |
|  66596 |
|  66961 |
+-----+
```

```
5 rows in set (0.30 sec)
```

```
mysql>
```

USE OF NOT EXISTS


To achieve the “**for all**” (universal quantifier- see Ch.8) effect, we use double negation this way in SQL:

Query: List first and last name of employees who work on ALL projects controlled by Dno=5.

```
SELECT Fname, Lname
FROM Employee
WHERE NOT EXISTS ( (SELECT Pnumber
                     FROM PROJECT
                     WHERE Dno=5)
```

Empty Set

```
      EXCEPT (SELECT Pno
                FROM WORKS_ON
                WHERE Ssn= ESsn));
```



The above is equivalent to double negation: List names of those employees for whom there does NOT exist a project managed by department no. 5 that they do NOT work on.

USE OF NOT EXISTS

```
mysql> SELECT salary FROM salaries WHERE NOT EXISTS (SELECT S.salary
from salaries AS S, current_dept_emp AS C WHERE S.emp_no = C.emp_no
AND C.dept_no = 'd001') limit 5;
Empty set (0.30 sec)
```

```
mysql>
```

```
mysql> SELECT salary FROM salaries WHERE salary NOT IN (SELECT
S.salary from salaries AS S, current_dept_emp AS C WHERE S.emp_no =
C.emp_no AND C.dept_no = 'd001') limit 5;
```

```
+-----+
| salary |
+-----+
|  81025 |
|  85112 |
|  69366 |
|  40006 |
|  43636 |
+-----+
```

Note that MySQL does not support the
EXCEPT, MINUS operator.

Double Negation to accomplish “for all” in SQL

- Q3B:

```
SELECT  Lname, Fname
FROM    EMPLOYEE
WHERE   NOT EXISTS ( SELECT *
                     FROM  WORKS_ON B
                     WHERE  ( B.Pno IN ( SELECT Pnumber
                                         FROM PROJECT
                                         WHERE Dnum=5)
                               AND
                               NOT EXISTS ( SELECT *
                                           FROM WORKS_ON C
                                           WHERE  C.Essn=Ssn
                                                  AND    C.Pno=B.Pno )));
```

Select any **WORK_ON(B)** tuples whose **Pno** is a project controlled by department 5

if there is not a **WORK_ON(C)** tuple with the same **Pno** and the same **Ssn**.

The above is a direct rendering of: List names of those employees for whom there does NOT exist a project managed by department no. 5 that they do NOT work on.

Explicit Sets and Renaming of Attributes in SQL

- Can use explicit set of values in WHERE clause

Q17: **SELECT DISTINCT Essn**
 FROM WORKS_ON
 WHERE Pno IN (1, 2, 3);

```
mysql> SELECT DISTINCT dept_no FROM current_dept_emp;
```

```
+-----+  
| dept_no |  
+-----+  
| d005    |  
| d007    |  
| d004    |  
| d003    |  
| d008    |  
| d006    |  
| d009    |  
| d001    |  
| d002    |  
+-----+
```

Explicit Sets and Renaming of Attributes in SQL

- Use qualifier **AS** followed by desired new name
 - Rename any attribute that appears in the result of a query

Q8A: **SELECT** E.Lname **AS** Employee_name, S.Lname **AS** Supervisor_name
 FROM EMPLOYEE **AS** E, EMPLOYEE **AS** S
 WHERE E.Super_ssn = S.Ssn;

```
mysql> SELECT C.dept_no AS department_number FROM
current_dept_emp AS C, employees AS E where C.emp_no =
E.emp_no limit 5;
```

```
+-----+
| department_number |
+-----+
| d005              |
| d007              |
| d004              |
| d004              |
| d003              |
+-----+
5 rows in set (0.29 sec)
```

Specifying Joined Tables in the FROM Clause of SQL

- **Joined table**
 - Permits users to specify a table resulting from a join operation in the FROM clause of a query
- **The FROM clause in Q1A**
 - Contains a single joined table. JOIN may also be called [INNER JOIN](#)

```
Q1A:  SELECT  Fname, Lname, Address
      FROM    (EMPLOYEE JOIN DEPARTMENT ON Dno = Dnumber)
      WHERE   Dname = 'Research';
```

Different Types of JOINed Tables in SQL

- Specify different types of join
 - NATURAL JOIN
 - Various types of OUTER JOIN (LEFT, RIGHT, FULL)
- **NATURAL JOIN on two relations R and S**
 - No join condition specified
 - Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

NATURAL JOIN

- Rename attributes of one relation so it can be joined with another using NATURAL JOIN:

```
Q1B:      SELECT      Fname, Lname, Address
           FROM        (EMPLOYEE NATURAL JOIN
                        (DEPARTMENT AS DEPT (Dname, Dno, Mssn,
                                             Msdate)))
           WHERE       Dname='Research';
```

The above works with EMPLOYEE.Dno = DEPT.Dno as an implicit join condition

NATURAL JOIN

```
mysql> SELECT C.dept_no AS department_number FROM current_dept_emp AS C  
NATURAL JOIN employees AS E limit 5;
```

department_number
d005
d007
d004
d004
d003


```
mysql> SELECT C.dept_no AS department_number FROM current_dept_emp AS C  
NATURAL JOIN employees AS E WHERE C.emp_no = E.emp_no limit 5;
```

department_number
d005
d007
d004
d004
d003

INNER and OUTER Joins

- **INNER JOIN (versus OUTER JOIN)**

- Default type of join in a joined table
- Tuple is included in the result only if a matching tuple exists in the other relation



- **LEFT OUTER JOIN**

- Every tuple in left table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of right table



- **RIGHT OUTER JOIN**

- Every tuple in right table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of left table



Example: LEFT OUTER JOIN

```
SELECT E.Lname AS Employee_Name  
       S.Lname AS Supervisor_Name  
FROM Employee AS E LEFT OUTER JOIN EMPLOYEE AS S  
                        ON E.Super_ssn = S.Ssn)
```

ALTERNATE SYNTAX:

```
SELECT E.Lname , S.Lname  
FROM   EMPLOYEE E, EMPLOYEE S  
WHERE  E.Super_ssn += S.Ssn
```


Example: LEFT OUTER JOIN

```
mysql> SELECT C.* FROM current_dept_emp AS C LEFT JOIN  
employees AS E ON C.dept_no = 'd001' limit 5;
```

emp_no	dept_no	from_date	to_date
10017	d001	1993-08-03	9999-01-01
10055	d001	1992-04-27	1995-07-22
10058	d001	1988-04-25	9999-01-01
10108	d001	1999-12-06	2001-10-20
10140	d001	1991-03-14	9999-01-01

5 rows in set (0.33 sec)

```
mysql>
```



Example: RIGHT OUTER JOIN

```
mysql> SELECT E.* FROM current_dept_emp AS C RIGHT JOIN employees AS E  
ON C.dept_no = 'd001' limit 5;
```

emp_no	birth_date	first_name	last_name	gender	hire_date
10001	1953-09-02	Georgi	Facello	M	1986-06-26
10001	1953-09-02	Georgi	Facello	M	1986-06-26
10001	1953-09-02	Georgi	Facello	M	1986-06-26
10001	1953-09-02	Georgi	Facello	M	1986-06-26
10001	1953-09-02	Georgi	Facello	M	1986-06-26

5 rows in set (0.29 sec)

```
mysql>
```




Example: CROSS JOIN (CARTESIAN PRODUCT)

```
mysql> SELECT E.* FROM current_dept_emp AS C CROSS JOIN employees AS E  
limit 5;
```

emp_no	birth_date	first_name	last_name	gender	hire_date
10001	1953-09-02	Georgi	Facello	M	1986-06-26
10002	1964-06-02	Bezalel	Simmel	F	1985-11-21
10003	1959-12-03	Parto	Bamford	M	1986-08-28
10004	1954-05-01	Chirstian	Koblick	M	1986-12-01
10005	1955-01-21	Kyoichi	Maliniak	M	1989-09-12

5 rows in set (0.30 sec)

```
mysql>
```



Multiway JOIN in the FROM clause

- **FULL OUTER JOIN** – combines result if **LEFT** and **RIGHT OUTER JOIN**
- Can nest JOIN specifications for a multiway join:

```
Q2A: SELECT  Pnumber, Dnum, Lname, Address, Bdate
          FROM      ((PROJECT JOIN DEPARTMENT ON
                      Dnum=Dnumber) JOIN EMPLOYEE ON
                      Mgr_ssn=Ssn)
          WHERE  Plocation='Stafford';
```

Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
 - COUNT, SUM, MAX, MIN, and AVG
- Grouping
 - Create subgroups of tuples before summarizing
- To select entire groups, HAVING clause is used
- Aggregate functions can be used in the SELECT clause or in a HAVING clause

Renaming Results of Aggregation

- **Following query returns a single row of computed values from EMPLOYEE table:**

```
Q19:      SELECT  SUM (Salary), MAX (Salary), MIN (Salary),
              AVG (Salary)
          FROM    EMPLOYEE;
```

- The result can be presented with new names:

```
Q19A:      SELECT      SUM (Salary) AS Total_Sal, MAX (Salary) AS
Highest_Sal, MIN (Salary) AS Lowest_Sal, AVG
(Salary) AS Average_Sal
FROM      EMPLOYEE;
```

Renaming Results of Aggregation

```
mysql> SELECT SUM(salary), AVG(salary), MIN(salary), MAX(salary)
from salaries;
```

SUM(salary)	AVG(salary)	MIN(salary)	MAX(salary)
181480757419	63810.7448	38623	158220

```
1 row in set (0.72 sec)
```

```
mysql>
```

Aggregate Functions in SQL (cont'd.)

- **NULL values are discarded when aggregate functions are applied to a particular column**

Query 20. Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:      SELECT      SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
           FROM        (EMPLOYEE JOIN DEPARTMENT ON Dno = Dnumber)
           WHERE        Dname = 'Research';
```

Queries 21 and 22. Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

```
Q21:      SELECT      COUNT (*)
           FROM        EMPLOYEE;
```

```
Q22:      SELECT      COUNT (*)
           FROM        EMPLOYEE, DEPARTMENT
           WHERE        DNO = DNUMBER AND DNAME = 'Research';
```


Aggregate Functions in SQL (cont'd.)

```
mysql> SELECT COUNT(*) from employees;
```

```
+-----+
```

```
| COUNT(*) |
```

```
+-----+
```

```
| 300024 |
```

```
+-----+
```

```
1 row in set (0.03 sec)
```

```
mysql>
```

Aggregate Functions on Booleans

- **SOME** and **ALL** may be applied as functions on Boolean Values.
- **SOME** returns true if at least one element in the collection is TRUE (similar to OR)
- **ALL** returns true if all of the elements in the collection are TRUE (similar to AND)

Grouping: The GROUP BY Clause

- **Partition relation into subsets of tuples**
 - Based on **grouping attribute(s)**
 - Apply function to each such group independently
- **GROUP BY clause**
 - Specifies grouping attributes
- **COUNT (*) counts the number of rows in the group**

Examples of GROUP BY

- The grouping attribute must appear in the SELECT clause:

Q24: SELECT Dno, COUNT (*), AVG (Salary)
 FROM EMPLOYEE
 GROUP BY Dno;

(a)

Fname	Minit	Lname	Ssn	...	Salary	Super_ssn	Dno
John	B	Smith	123456789		30000	333445555	5
Franklin	T	Wong	333445555		40000	888665555	5
Ramesh	K	Narayan	666884444		38000	333445555	5
Joyce	A	English	453453453	...	25000	333445555	5
Alicia	J	Zelaya	999887777		25000	987654321	4
Jennifer	S	Wallace	987654321		43000	888665555	4
Ahmad	V	Jabbar	987987987		25000	987654321	4
James	E	Bong	888665555		55000	NULL	1

Dno	Count (*)	Avg (Salary)
5	4	33250
4	3	31000
1	1	55000

Result of Q24

Grouping EMPLOYEE tuples by the value of Dno

```
mysql> SELECT dept_no, COUNT(*)
FROM current_dept_emp GROUP BY
dept_no;
```

dept_no	COUNT(*)
d001	18426
d002	15579
d003	16071
d004	66675
d005	76958
d006	18295
d007	46922
d008	19285
d009	21813

Examples of GROUP BY

- If the grouping attribute has NULL as a possible value, then a separate group is created for the null value (e.g., null Dno in the above query)
- GROUP BY may be applied to the result of a JOIN:
Q25: SELECT Pnumber, Pname, COUNT (*)
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname;

Examples of GROUP BY

```
mysql> SELECT dept_no, gender, COUNT(*) FROM current_dept_emp C,  
employees E WHERE C.emp_no = E.emp_no GROUP BY dept_no, E.gender;
```

dept_no	gender	COUNT(*)
d001	M	11111
d001	F	7315
d002	M	9273
d002	F	6306
d003	M	9701
d003	F	6370
d004	M	39885
d004	F	26790
d005	M	46218
d005	F	30740
d006	M	10921
d006	F	7374
d007	M	28176
d007	F	18746
d008	M	11587
d008	F	7698
d009	M	13101
d009	F	8712

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

- **HAVING** clause
 - Provides a condition to select or reject an entire group:
- Query 26. For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

Q26:	SELECT	Pnumber, Pname, COUNT (*)
	FROM	PROJECT, WORKS_ON
	WHERE	Pnumber=Pno
	GROUP BY	Pnumber, Pname
	HAVING	COUNT (*) > 2;

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

(b)

Pname	Pnumber	...	Essn	Pno	Hours
ProductX	1		123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666884444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10	...	333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by the HAVING condition of Q26.

After applying the WHERE clause but before applying HAVING

Pname	Pnumber	...	Essn	Pno	Hours
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
Computerization	10		333445555	10	10.0
Computerization	10	...	999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

Pname	Count (*)
ProductY	3
Computerization	3
Reorganization	3
Newbenefits	3

Result of Q26
(Pnumber not shown)

After applying the HAVING clause condition

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

```
mysql> SELECT dept_no, gender, COUNT(*) FROM current_dept_emp C,  
employees E WHERE C.emp_no = E.emp_no GROUP BY dept_no, E.gender  
HAVING COUNT(*) > 10000;
```

dept_no	gender	COUNT(*)
d001	M	11111
d004	M	39885
d004	F	26790
d005	M	46218
d005	F	30740
d006	M	10921
d007	M	28176
d007	F	18746
d008	M	11587
d009	M	13101

```
+-----+-----+-----+  
10 rows in set (1.07 sec)
```

```
mysql>
```

Combining the WHERE and the HAVING Clause

- Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.

- INCORRECT QUERY:**

```
SELECT      Dno, COUNT (*)  
FROM        EMPLOYEE  
WHERE       Salary > 40000  
GROUP BY   Dno  
HAVING      COUNT (*) > 5;
```

it will select only departments that have more than five employees who each earn more than \$40000.

The rule is that the **WHERE** clause is executed first, to select individual tuples or joined tuples; the **HAVING** clause is applied later, to select individual groups of tuples.

Combining the WHERE and the HAVING Clause (continued)

Correct Specification of the Query:

- **Note:** the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

```
Q28:  SELECT      Dno, COUNT (*)
        FROM      EMPLOYEE
        WHERE      Salary>40000 AND Dno IN
                   ( SELECT      Dno
                     FROM      EMPLOYEE
                     GROUP BY    Dno
                     HAVING      COUNT (*) > 5)
        GROUP BY    Dno;
```

Use of WITH

- The WITH clause allows a user to define a table that will only be used in a particular query (**not available in all SQL implementations**)
- Used for convenience to create a temporary “View” and use that immediately in a query
- Allows a more straightforward way of looking a step-by-step query

A **common table expression (CTE)** is a named temporary result set that exists within the scope of a single statement and that can be referred to later within that statement, possibly multiple times.

Example of WITH

- See an alternate approach to doing Q28:

```
Q28':  WITH BIGDEPTS (Dno) AS
        (SELECT Dno
         FROM EMPLOYEE
         GROUP BY Dno
         HAVING      COUNT (*) > 5)

SELECT  Dno, COUNT (*)
FROM    EMPLOYEE
WHERE   Salary > 40000 AND Dno IN BIGDEPTS
GROUP BY Dno;
```

Use of CASE

- **SQL also has a CASE construct**
- **Used when a value can be different based on certain conditions.**
- **Can be used in any part of an SQL query where a value is expected**
- **Applicable when querying, inserting or updating tuples**

EXAMPLE of use of CASE

- The following example shows that employees are receiving different raises in different departments (A variation of the update U6)

U6':

UPDATE	EMPLOYEE		
SET	Salary =		
CASE	WHEN Dno = 5	THEN	Salary + 2000
	WHEN Dno = 4	THEN	Salary + 1500
	WHEN Dno = 1	THEN	Salary + 3000

EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```


Specifying Constraints as Assertions and Actions as Triggers

Specifying Constraints as Assertions and Actions as Triggers

- **Semantic Constraints:** The following are beyond the scope of the EER and relational model
- **CREATE ASSERTION**
 - Specify additional types of constraints outside scope of built-in relational model constraints
- **CREATE TRIGGER**
 - Specify automatic actions that database system will perform when certain events and conditions occur

Specifying General Constraints as Assertions in SQL

- **CREATE ASSERTION**

- Specify a query that selects any tuples that violate the desired condition
- Use only in cases where it goes beyond a simple **CHECK** which applies to individual attributes and domains

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT      *
                      FROM        EMPLOYEE E, EMPLOYEE M,
                                DEPARTMENT D
                      WHERE       E.Salary > M.Salary
                                AND      E.Dno = D.Dnumber
                                AND      D.Mgr_ssn = M.Ssn ) );
```

Introduction to Triggers in SQL

- **CREATE TRIGGER** statement
 - Used to monitor the database
- Typical trigger has three components which make it a rule for an “active database”:
 - Event(s)
 - Condition
 - Action

USE OF TRIGGERS

- **AN EXAMPLE with standard Syntax.(Note : other SQL implementations like PostgreSQL use a different syntax.)**

R5:

**CREATE TRIGGER SALARY_VIOLATION
BEFORE INSERT OR UPDATE OF Salary, Supervisor_ssn ON
EMPLOYEE**

FOR EACH ROW

**WHEN (NEW.SALARY > (SELECT Salary FROM EMPLOYEE
WHERE Ssn = NEW. Supervisor_Ssn))
INFORM_SUPERVISOR (NEW.Supervisor.Ssn, New.Ssn)**

Views (Virtual Tables) in SQL

- **Concept of a view in SQL**
 - Single table derived from other tables called the **defining tables**
 - Considered to be [a virtual table that is not necessarily populated](#)

Specification of Views in SQL

- **CREATE VIEW** command

- Give table name, list of attribute names, and a query to specify the contents of the view
- In V1, attributes retain the names from base tables. In V2, attributes are assigned names

V1:	CREATE VIEW	WORKS_ON1
	AS SELECT	Fname, Lname, Pname, Hours
	FROM	EMPLOYEE, PROJECT, WORKS_ON
	WHERE	Ssn = Essn AND Pno = Pnumber;
V2:	CREATE VIEW	DEPT_INFO(Dept_name, No_of_emps, Total_sal)
	AS SELECT	Dname, COUNT (*), SUM (Salary)
	FROM	DEPARTMENT, EMPLOYEE
	WHERE	Dnumber = Dno
	GROUP BY	Dname;

Specification of Views in SQL

V1: **CREATE VIEW** WORKS_ON1
 AS SELECT Fname, Lname, Pname, Hours
 FROM EMPLOYEE, PROJECT, WORKS_ON
 WHERE Ssn = Essn **AND** Pno = Pnumber;

V2: **CREATE VIEW** DEPT_INFO(Dept_name, No_of_emps, Total_sal)
 AS SELECT Dname, **COUNT** (*), **SUM** (Salary)
 FROM DEPARTMENT, EMPLOYEE
 WHERE Dnumber = Dno
 GROUP BY Dname;

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------

DEPT_INFO

Dept_name	No_of_emps	Total_sal
-----------	------------	-----------

Specification of Views in SQL (cont'd.)

- Once a View is defined, SQL queries can use the View relation in the FROM clause
- View is always up-to-date
 - Responsibility of the DBMS and not the user
- **DROP VIEW** command
 - Dispose of a view

View Implementation, View Update, and Inline Views

- **Complex problem of efficiently implementing a view for querying**
- **Strategy1: Query modification approach**
 - Compute the view as and when needed. Do not store permanently
 - Modify view query into a query on underlying base tables
 - Disadvantage: inefficient for views defined via complex queries that are time-consuming to execute

View Materialization

- **Strategy 2: View materialization**
 - Physically create a temporary view table when the view is first queried
 - Keep that table on the assumption that other queries on the view will follow
 - Requires efficient strategy for automatically updating the view table when the base tables are updated
- **Incremental update strategy for materialized views**
 - DBMS determines what new tuples must be inserted, deleted, or modified in a materialized view table

View Materialization (contd.)

- **Multiple ways to handle materialization:**
 - **immediate update** strategy updates a view as soon as the base tables are changed
 - **lazy update** strategy updates the view when needed by a view query
 - **periodic update** strategy updates the view periodically (in the latter strategy, a view query may get a result that is not up-to-date). This is commonly used in Banks, Retail store operations, etc.

View Update

- **Update on a view defined on a single table without any aggregate functions**
 - Can be mapped to an update on underlying base table- possible if the primary key is preserved in the view
- **Update not permitted on aggregate views. E.g.,**

UV2:

UPDATE	DEPT_INFO
SET	Total_sal=100000
WHERE	Dname='Research';

cannot be processed because Total_sal is a computed value in the view definition

View Update

- In MySQL, views are not only query-able but also updatable. It means that you can use the INSERT or UPDATE statement to insert or update rows of the base table through the updatable view. In addition, you can use DELETE statement to remove rows of the underlying table through the view.
- However, to create an updatable view, the SELECT statement that defines the view must not contain any of the following elements:
 - Aggregate functions such as MIN, MAX, SUM, AVG, and COUNT.
 - DISTINCT
 - GROUP BY clause.
 - HAVING clause.
 - UNION or UNION ALL clause.
 - Left join or outer join.
 - Subquery in the SELECT clause or in the WHERE clause that refers to the table appeared in the FROM clause.
 - Reference to non-updatable view in the FROM clause.
 - Reference only to literal values.
 - Multiple references to any column of the base table.

View Update and Inline Views

- **View involving joins**
 - Often not possible for DBMS to determine which of the updates is intended
- **Clause WITH CHECK OPTION**
 - Must be added at the end of the view definition if a view is to be updated to make sure that tuples being updated stay in the view
- **In-line view**
 - Defined in the FROM clause of an SQL query (e.g., we saw its used in the WITH example)

Views as authorization mechanism

- SQL query authorization statements (GRANT and REVOKE) are described in detail in Chapter 30
- Views can be used to hide certain attributes or tuples from unauthorized users
- E.g., For a user who is only allowed to see employee information for those who work for department 5, he may only access the view **DEPT5EMP**:

```
CREATE VIEW      DEPT5EMP  AS
SELECT          *
FROM            EMPLOYEE
WHERE           Dno = 5;
```


Views as authorization mechanism

```
mysql> CREATE VIEW TEST AS SELECT * FROM dept_emp WHERE  
dept_no = 'd001';  
Query OK, 0 rows affected (0.69 sec)
```

```
mysql>
```

```
mysql> SELECT * FROM TEST limit 5;
```

emp_no	dept_no	from_date	to_date
10017	d001	1993-08-03	9999-01-01
10055	d001	1992-04-27	1995-07-22
10058	d001	1988-04-25	9999-01-01
10108	d001	1999-12-06	2001-10-20
10140	d001	1991-03-14	9999-01-01

```
5 rows in set (0.00 sec)
```

```
mysql> DROP VIEW TEST;
```

```
Query OK, 0 rows affected (0.00 sec)
```

Schema Change Statements in SQL

Schema Change Statements in SQL

- **Schema evolution commands**
 - DBA may want to change the schema while the database is operational
 - Does not require recompilation of the database schema

The DROP Command

- **DROP command**
 - Used to drop named schema elements, such as tables, domains, or constraint
- **Drop behavior options:**
 - `CASCADE` and `RESTRICT`
- **Example:**
 - `DROP SCHEMA COMPANY CASCADE;`
 - This removes the schema and all its elements including tables, views, constraints, etc.

The ALTER table command

- **Alter table actions include:**
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints
- **Example:**
 - `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12);`

Adding and Dropping Constraints

- **Change constraints specified on a table**
 - Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```

Dropping Columns, Default Values

- **To drop a column**
 - Choose either **CASCADE** or **RESTRICT**
 - **CASCADE** would drop the column from views etc. **RESTRICT** is possible if no views refer to it.

```
ALTER TABLE COMPANY.EMPLOYEE  
    DROP COLUMN Address CASCADE;
```

- **Default values can be dropped and altered :**

```
ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn  
    DROP DEFAULT;
```

```
ALTER TABLE COMPANY.DEPARTMENT  
    ALTER COLUMN Mgr_ssn SET DEFAULT '333445555';
```

Table 7.2 Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

```
CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]  
                             { , <column name> <column type> [ <attribute constraint> ] }  
                             [ <table constraint> { , <table constraint> } ] )
```

```
DROP TABLE <table name>  
ALTER TABLE <table name> ADD <column name> <column type>
```

```
SELECT [ DISTINCT ] <attribute list>  
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }  
[ WHERE <condition> ]  
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]  
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) )  
                    { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) } ) )
```

```
<grouping attributes> ::= <column name> { , <column name> }
```

```
<order> ::= ( ASC | DESC )
```

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]  
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) } )  
| <select statement> )
```

continued on next slide

Table 7.2 (continued) Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

DELETE FROM <table name>

[WHERE <selection condition>]

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[WHERE <selection condition>]

CREATE [UNIQUE] INDEX <index name>

ON <table name> (<column name> [<order>] { , <column name> [<order>] })

[CLUSTER]

DROP INDEX <index name>

CREATE VIEW <view name> [(<column name> { , <column name> })]

AS <select statement>

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

Summary

- **Complex SQL:**
 - Nested queries, joined tables (in the FROM clause), outer joins, aggregate functions, grouping
- **Handling semantic constraints with `CREATE ASSERTION` and `CREATE TRIGGER`**
- **`CREATE VIEW` statement and materialization strategies**
- **Schema Modification for the DBAs using `ALTER TABLE`, `ADD` and `DROP COLUMN`, `ALTER CONSTRAINT` etc.**